

Introduction to Reinforcement Learning

Mattia Pellegrino, Ph.D Fellow
mattia.pellegrino@unipr.it



Summary


- Information about the course
- About Reinforcement Learning
- The Reinforcement Learning challenges
- RL Agent
- Reinforcement Learning Problems

COURSE INFORMATION

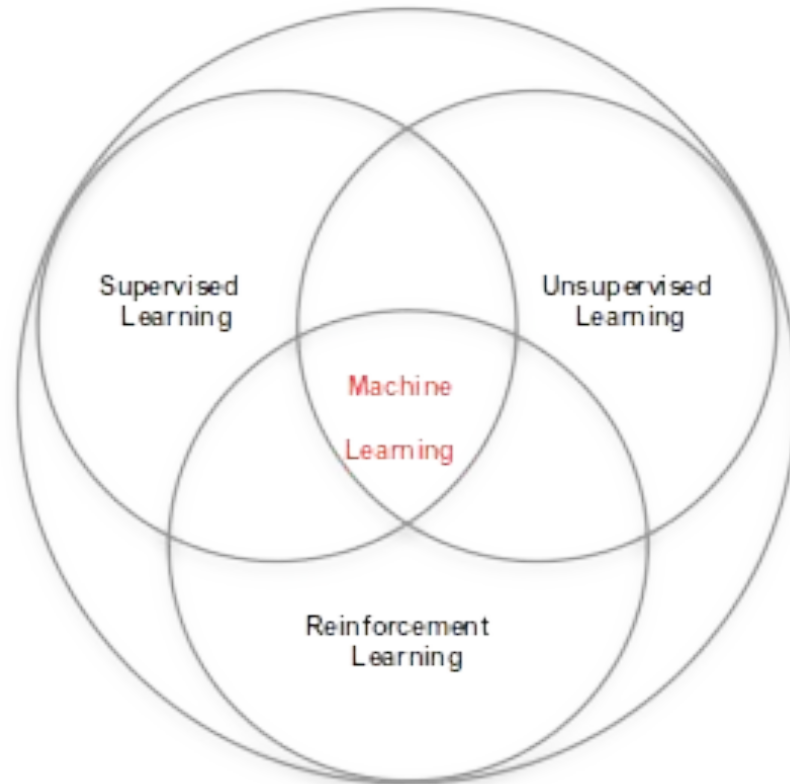
- Contact me: mattia.pellegrino@unipr.it
- **TextBooks:**
 - Montague, P. Read. "Reinforcement learning: an introduction, by Sutton, RS and Barto, AG." Trends in cognitive sciences 3.9 (1999): 360. <http://webdocs.cs.ualberta.ca/~sutton/book/the-book.html>
 - Appress, "**Deep Reinforcement Learning with python**, by Nimish Sanghi", (2021)

RL PLACEMENT

- All these branches try to solve the same problem: “**Decision Making**”
- Find the best actions’ combination to get the best result




MACHINE LEARNING BRANCHES



RL: PROPERTIES

- What makes RL different from other learning methodologies?
 - There is no supervisor. Instead, there is a *reward signal*
 - *How do we know if what we are doing is right or not?*
 - Reward is not instantaneous
 - *We realize that we have done a bad choice afterwards*
 - Time matters
 - Agent's actions affect the entire environment

RL: A ROBOT LEARNS TO WALK




RL: TRACKMANIA



RL: BOXING



RL: ATARI GAME



RL: REWARD

- A reward R_t is a scalar feedback signal
 - It is just a number
- Indicates how good is the action choose by the agent at **timestep t**
- The agent's goal is to maximize the total reward

RL is based on the ***reward hypothesis***

“All goals can be described by the maximization of expected cumulative reward”

RL: EXAMPLE OF REWARDS


- **A robot learns to walk**
 - + reward for forward motion
 - - reward for falling over
- **Trackmania**
 - + reward for forward motion
 - - reward to falling over
- **Boxing**
 - + reward to stand stand in place
 - - reward to falling over
 - + reward to hit the enemy
 - - reward to get hit
- **Play Atari Games**
 - +/- reward for increasing/decreasing score

RL: SEQUENTIAL DECISION MAKING


- **Goal:** select actions to maximize total future reward
- Actions may have long term consequences
- Reward may be **delayed**
- It may be better to **sacrifice immediate reward** to gain more **long-term reward**
- Examples:
 - A financial investment (may take months to mature)
 - Blocking opponent moves (might help winning chances many moves from now)

AGENT AND ENVIRONMENT

- We will use the current formalism



AGENT AND ENVIRONMENT



- At each step t the agent:
 - Execute an action A_t
 - Obtain an observation O_t
 - Obtain a reward R_t
- The environment
 - Receives an action A_t
 - Releases an observation O_{t+1}
 - Releases a Reward R_{t+1}

HISTORY - STATE

History is the sequence of observation, action, and rewards (usually huge)

$$H_t = O_1, R_1, A_1, \dots, A_t, O_t, R_t$$


- An algorithm is a **mapping** between the **history** and what **happens next**:
 - The agent selects actions
 - The environment selects observations/rewards

The state is the information used to determine what happens next

- It's a **synthesis** of what happened and we base on this because the history usually is too big to compute
- Formally, state is a function of the history


$$S_t = f(H_t)$$

ENVIRONMENT STATE



- The ***environment state*** S_t^e is the environment's private representation
- This state is **not** usually **visible** to the agent
- Sometimes contains **irrelevant** information

AGENT STATE



- The **agent state** S_t^a is the agent's internal representation
- Whatever information the agent can use to choose the next action
- It can be any function of history:

$$S_t = f(H_t)$$

INFORMATION STATE

An information state (a.k.a. Markov state) contains all useful information from the history

- *A state S_t is a Markov state if and only if*


$$\mathbb{P}[S_{t+1} \mid S_t] = \mathbb{P}[S_{t+1} \mid S_1, \dots, S_t]$$

- The **future** is **independent** of the **past** given present

$$H_{1:t} \longrightarrow S_t \longrightarrow H_{t+1:\infty}$$

- Once the **state** is **known**, the **history** is **irrelevant**
- The environment state S_t^e is a Markov State
- The history H_t is a Markov state

FULLY OBSERVABLE ENVIRONMENTS



- Full observability: agent directly observes environment state

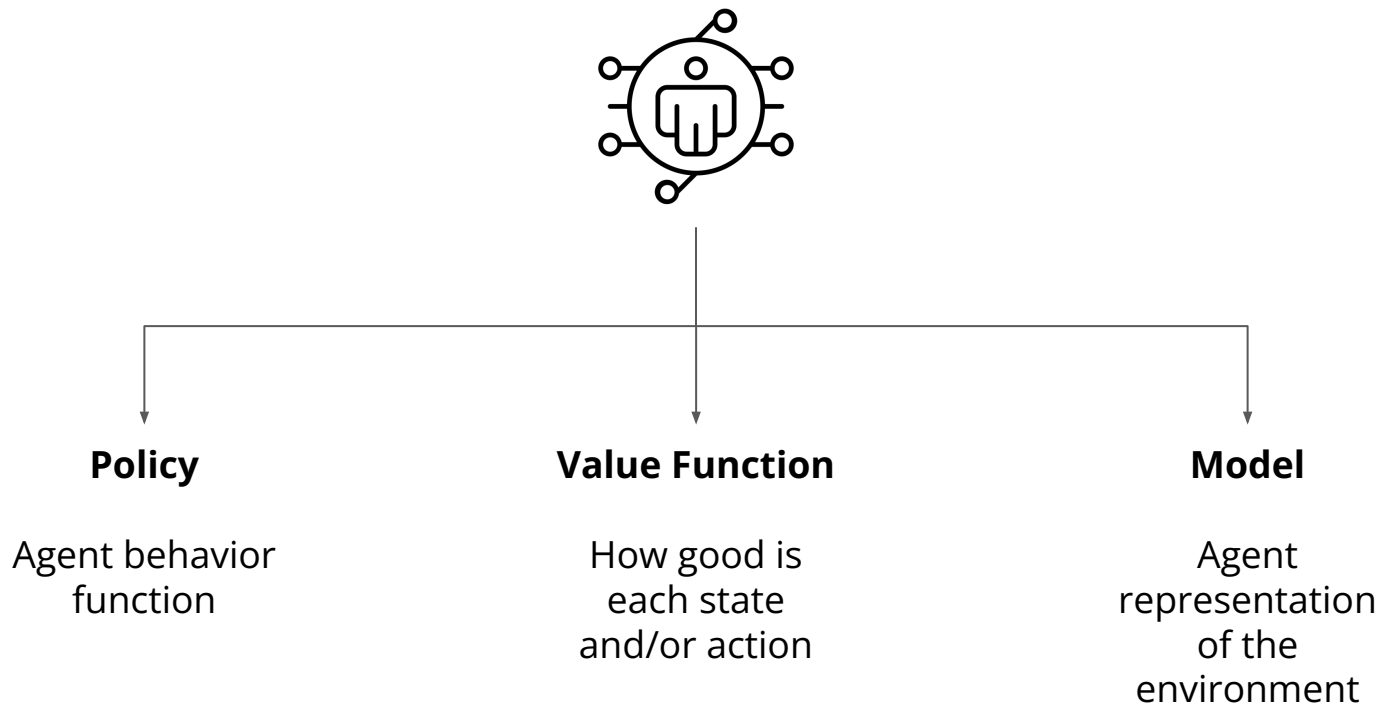
$$O_t = S_t^a = S_t^e$$

- Agent state = environment state = information state
- Formally, this is a Markov Decision Process (MDP)

PARTIALLY OBSERVABLE ENVIRONMENTS

- **Partial observability:** agent indirectly observes environment:
 - A robot with a camera vision
 - A card player agent
 - A trading agents
- In this case, the agent and the environment state are not the same
- This is a **partially observable Markov decision process (POMDP)**
- An agent must build its own state representation (S_a^t)

MAJOR RL AGENT COMPONENTS



POLICY

- A **policy** is the agent's behaviour
- It is a map from state of action
- There are 2 types of policies:
 - Deterministic policy: $a = \pi(s)$
 - Stochastic policy: $\pi(a \mid s) = \mathbb{P}[A_t = a \mid S_t = s]$

VALUE FUNCTION

- Value function is a prediction of future reward
- We use it to evaluate the goodness of the various states
- We can use it to pick the best action

$$v_{\pi}(s) = \mathbb{E}_{\pi} [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$$


MODEL

- A model predicts what the environment will do next
- P predicts the next state
- R predicts the next (immediate) rewards

$$P_{ss'}^a = P[S_{t+1} = s' \mid S_t = s, A_t = a]$$


$$R_s^a = E[R_{t+1} \mid S_t = s, A_t = a]$$

MAZE EXAMPLE




- Rewards: -1 per time-step
- Action: L,R,U,D
- States: Agent's Location

MAZE EXAMPLE



- Arrows represent policy $\pi(s)$ for each state s

MAZE EXAMPLE



Numbers represent value $v_{\pi}(s)$ of each state s

EXPLORATION AND EXPLOITATION

- Reinforcement learning is like trial-and-error learning
- The agents should discover a good policy from its experiences of the environment
- Without losing too much reward along the way
- Exploration find more information about the environment
- Exploitation exploits known information to maximize reward

RL AGENT TAXONOMY

